

# 1 Linear Regression using Gradient Descent

## Coding in Python (75 points)

For this part, you will write your own code in Python for implementing the gradient descent algorithm, and apply it to a linear regression problem. You are free to use any data loading, pre-processing, parsing, and graphing library, such as numpy, pandas, graphics. **However, you cannot use any library that implements gradient descent or linear regression.**

You will need to perform the following:

1. Choose a dataset suitable for regression from UCI ML Repository: - <https://archive.ics.uci.edu/ml/datasets.php>. If the above link doesn't work, you can go to the main page: <https://archive.ics.uci.edu/ml/index.php> and choose "view all datasets option". Host the dataset on a public location e.g. UTD web account. Please do not hard code paths to your local computer.
2. Pre-process your dataset. Pre-processing includes the following activities:
  - Remove null or NA values
  - Remove any redundant rows
  - Convert categorical variables to numerical variables
  - If you feel an attribute is not suitable or is not correlated with the outcome, you might want to get rid of it.
  - Any other pre-processing that you may need to perform.
3. After pre-processing split the dataset into training and test parts. It is up to you to choose the train/test ratio, but commonly used values are 80/20, 90/10, etc.
4. Use the training dataset to construct a linear regression model. We discussed creating a linear regression model using a single attribute in class. For this assignment, you will need to extend that model to consider multiple attributes. You may want to think of the vector form of the weight update equation. **Note again: you cannot use a library that implements gradient descent or linear regression.**

There are various parameters such as *learning rate*, *number of iterations* or other *stopping condition*, etc. You need to tune these parameters to achieve the optimum error value. Tuning involves testing various combinations, and then using the best one. You need to create a log file that indicates parameters used and error (MSE) value obtained for various trials.

5. Apply the model you created in the previous step to the test part of the dataset. Report the test dataset error values for the best set of parameters obtained from previous part. If you are not satisfied with your answer, you can repeat the training step.
6. Answer this question: Are you satisfied that you have found the best solution? Explain.

## 2 Linear Regression using ML libraries (25 points)

In the second part of this assignment, you will use any ML library that performs linear regression from Scikit Learn package. <https://scikit-learn.org> on the **same dataset that you used in part 1**.

Use similar workflow like in part 1. The difference would be that the package will build the model for you. Report output similar to earlier part.

## 3 Additional Requirements

- You will be graded on the basis of your code. Write clean and elegant code that should be in the form of Python classes, with appropriate constructors, and other methods.
- Parameters used such as number of iterations, learning rate, should be optimized. Keep a log file of your trials indicating parameters used, training error, and test error.
- You need to provide as many plots as possible. They could be MSE vs number of iterations, the output variable plotted against one or more of important attributes. Use your judgement and be creative.
- Output as many evaluation statistics as possible. Some examples are weight coefficients, MSE,  $R^2$  value, [https://en.wikipedia.org/wiki/Coefficient\\_of\\_determination](https://en.wikipedia.org/wiki/Coefficient_of_determination), Explained Variance [https://en.wikipedia.org/wiki/Explained\\_variation](https://en.wikipedia.org/wiki/Explained_variation), and any other
- Be sure to answer the following question: Are you satisfied that the package has found the best solution. How can you check. Explain.

## 4 What to Submit

- For parts 1 and 2, completed Python code file(s). Remember to properly name your files such as part1.py and part2.py.

- README file indicating how to build and run your code. For part 2, indicate which libraries you have used. If the TA cannot run your code, you don't get any credit.
- **Please do not submit any dataset files.** Host your data on a public web source, such as your UTD web account, AWS, etc.
- Do not hardcode paths on your local computer.
- A report file containing log of your trials with different parameters, answer to questions, and plots.